

Chimpanzees, part 2

Stat 341 — Spring 2017

April, 2017

The Data

```
library(rethinking)
data(chimpanzees)
Chimps <- chimpanzees %>%
  mutate(
    recipient = ifelse(is.na(recipient), -1, recipient), # to avoid missing data
    combo = paste0(prosoc_left, "/", condition)         # useful for plotting
  )
Chimps %>% sample(3) # 3 random rows
```

	actor	recipient	condition	block	trial	prosoc_left	chose_prosoc	pulled_left	combo	orig.id
## 225	4	-1	0	2	17	0	0	1	0/0	225
## 365	6	-1	0	1	10	0	0	1	0/0	365
## 96	2	-1	0	4	47	1	1	1	1/0	96

Model with main effects, interaction, and varying intercepts

Note: This is not the same as model m12.4 in the text.

```
m12.4a <- map2stan(
  alist(
    pulled_left ~ dbinom(1, p),
    logit(p) <- a + a_actor[actor] +
      bp * prosoc_left + bc * condition + bpc * prosoc_left * condition,
    a_actor[actor] ~ dnorm(0, sigma_actor),
    c(a, bp, bc, bpc) ~ dnorm(0, 10), # nifty short cut!
    sigma_actor ~ dcauchy(0, 1)
  ),
  data = Chimps,
  warmup = 1000, iter = 5000, chains = 4, cores = 3, refresh = 0
)
```

1. Describe how to modify the code above in a way that removes `a` from the second line of the formula list but still fits the “same model”.
2. In what sense are the two models “the same”? In what ways are they different?


Posterior samples

```
m12.4a_post <- extract.samples(m12.4a)
glimpse(m12.4a_post)
```

```
## List of 6
## $ a_actor      : num [1:16000, 1:7] -0.9048 -1.0586 0.3696 -0.5498 0.0517 ...
## $ a            : num [1:16000(1d)] 0.1379 0.3277 -0.8717 0.0158 0.0439 ...
## $ bp           : num [1:16000(1d)] 0.647 0.916 0.688 0.335 0.329 ...
## $ bc           : num [1:16000(1d)] -0.505 -0.287 -0.319 -0.382 -0.658 ...
## $ bpc          : num [1:16000(1d)] 0.5761 0.0208 0.3456 0.4342 1.0619 ...
## $ sigma_actor : num [1:16000(1d)] 1.69 1.54 2.39 1.65 4.14 ...
```

3. How does the result of `extract.samples()` change now that we have a multi-level model?
4. In a varying intercepts model, each actor has its own intercept. In this model, what does the intercept for actor 7 represent?
5. In a Bayesian model, there is a posterior distribution for this intercept. Write a little R code that plots the posterior distribution of the intercept for actor 7.

```
gf_dens(~intercept.7, data = m12.4a_postD)
gf_dens(~logistic(intercept.7), data = m12.4a_postD)
```



A model with multiple clusters

Again, this is not quite the same as in the book. I've added the other main effect.

```
# prep data
Chimps <-
  Chimps %>% rename(block_id = block) # name 'block' is reserved by Stan


m12.5a <- map2stan(
  alist(
    pulled_left ~ dbinom(1, p),
    logit(p) <- a + a_actor[actor] + a_block[block_id] +
      bp * prosoc_left + bc * condition + bpc * condition * prosoc_left,
    a_actor[actor] ~ dnorm(0, sigma_actor),
    a_block[block_id] ~ dnorm(0, sigma_block),
    c(a, bp, bc, bpc) ~ dnorm(0, 10), # short cut again
    c(sigma_actor, sigma_block) ~ dcauchy(0, 1) # and again
  ),
  data = Chimps, warmup = 1000, iter = 6000, chains = 4, cores = 3, refresh = 0
)
```

6. Explain how `m12.5a` differs from `m12.4a`.

```
coef(m12.5a)
```

```
## a_actor[1] a_actor[2] a_actor[3] a_actor[4] a_actor[5] a_actor[6] a_actor[7]
## -1.19545192 4.24205314 -1.49905491 -1.50601607 -1.19542462 -0.23643461 1.31639119
## a_block[1] a_block[2] a_block[3] a_block[4] a_block[5] a_block[6] a
## -0.17768577 0.03651059 0.05526852 0.01376358 -0.03854017 0.12089453 0.67240324
## bp bc bpc sigma_actor sigma_block
## 0.62792839 -0.41643357 0.28225948 2.32175193 0.22768948
```

```
# precis(m12.5a, depth = 2) # saving some paper
plot(precis(m12.5a, depth = 2)) # just show the plot
```



7. What do you notice when you compare actors to blocks?

8. How do the comparisons below fit into this story?

```
compare(m12.4a, m12.5a)@output %>%
  mutate(
    logLik = c(logLik(m12.4a), m12.5a = logLik(m12.5a)),
    dlogLik = logLik - logLik(m12.4a))
```

```
##      WAIC    pWAIC    dWAIC   weight      SE      dSE   logLik dlogLik
## 1 531.7753  9.21556 0.000000 0.6315959 19.76806      NA -256.1056 0.000000
## 2 532.8534 11.43523 1.078138 0.3684041 20.00556 1.783302 -254.3809 1.724775
```

Posterior samples

Here's a quick trick for turning the list produced by `extract.samples()` into a data frame you can use for plotting. Notice the way the names change.

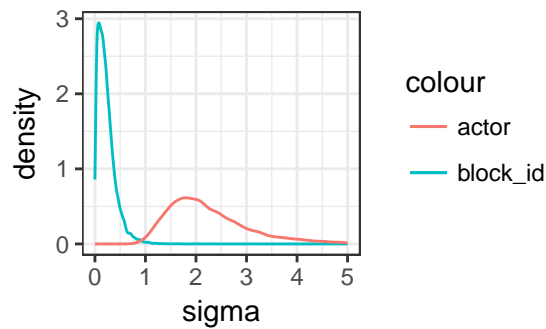
```
m12.5a_post <- extract.samples(m12.5a) %>% data.frame()
names(m12.5a_post)
```

```
## [1] "a_actor.1" "a_actor.2" "a_actor.3" "a_actor.4" "a_actor.5" "a_actor.6"
```

```
## [7] "a_actor.7"    "a_block.1"    "a_block.2"    "a_block.3"    "a_block.4"    "a_block.5"
## [13] "a_block.6"    "a"            "bp"           "bc"           "bpc"          "sigma_actor"
## [19] "sigma_block"
```

```
gf_dens(~sigma_block + color::"block_id", data = m12.5a_post) %>%
  gf_dens(~sigma_actor + color::"actor", data = m12.5a_post) %>%
  gf_labs(x = "sigma") %>%
  gf_lims(x = c(0,5))
```

```
## Warning: Removed 445 rows containing non-finite values (stat_density).
```



Converting to a “long” format.

If we convert our posterior samples from a “wide” format (one row per sample) to a “long” format (one row for each sample of each parameter), we can plot posterior distributions for multiple parameters even more easily.

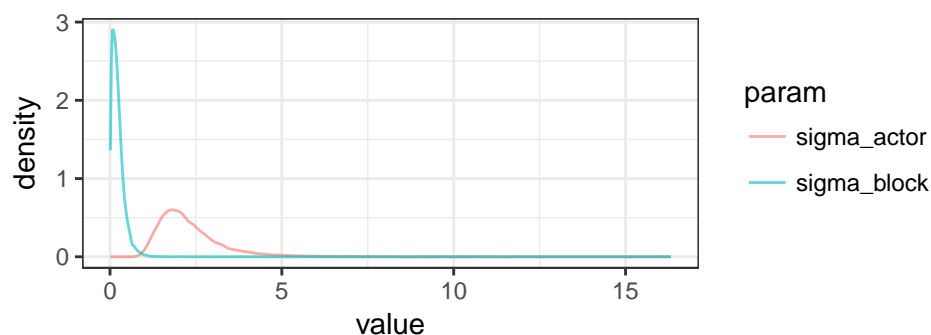
```
m12.5a_post_long <- m12.5a_post %>%
  tidyr::gather(param, value)
```

```
## Warning: attributes are not identical across measure variables; they will be dropped
```

```
m12.5a_post_long %>% head(3)
```

```
##      param      value
## 1 a_actor.1 -0.7378205
## 2 a_actor.1 -0.5102513
## 3 a_actor.1  0.8929754
```

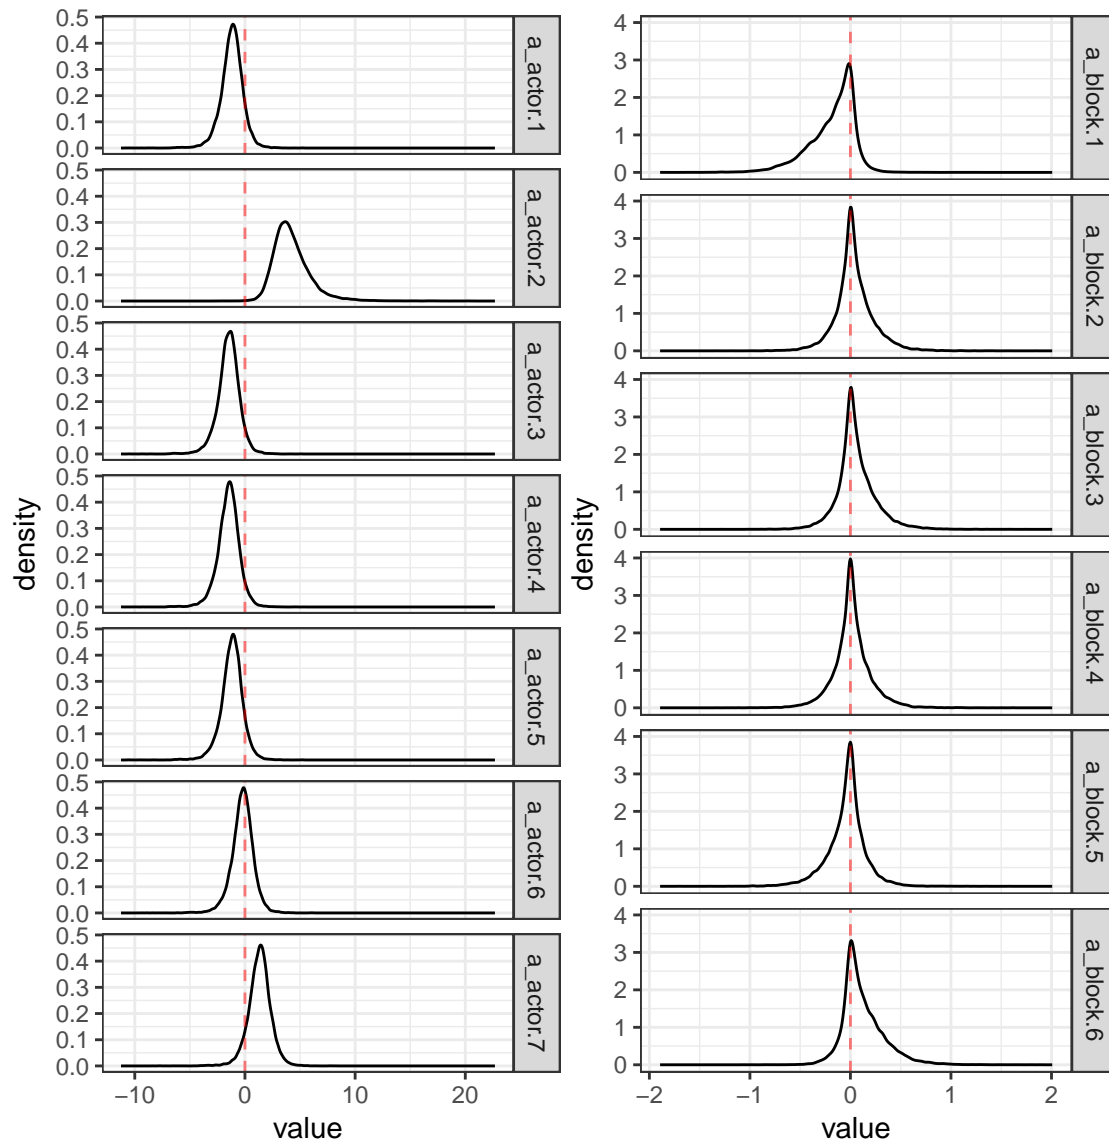
```
gf_dens(~ value + color::param, alpha = 0.6,
  data = m12.5a_post_long %>% filter(grepl("sigma", param)))
```



```

gf_dens(~ value,
  data = m12.5a_post_long %>% filter(grepl("a_actor\\.", param))) %>%
  gf_vline(xintercept = 0, color = "red", alpha = 0.5, linetype = "dashed") %>%
  gf_facet_grid( param ~ .)
gf_dens(~ value,
  data = m12.5a_post_long %>% filter(grepl("a_block\\.", param))) %>%
  gf_vline(xintercept = 0, color = "red", alpha = 0.5, linetype = "dashed") %>%
  gf_facet_grid( param ~ .)

```



R Code 12.27

```
m12.4a_pred <-
  expand.grid(
    prosoc_left = 0:1,
    condition = 0:1,
    actor = 1:7) %>%
  mutate(
    combo = paste0(prosoc_left, "/", condition)
  )


link.m12.4a <- link(m12.4a, data = m12.4a_pred)
```

```
## [ 100 / 1000 ]
[ 200 / 1000 ]
[ 300 / 1000 ]
[ 400 / 1000 ]
[ 500 / 1000 ]
[ 600 / 1000 ]
[ 700 / 1000 ]
[ 800 / 1000 ]
[ 900 / 1000 ]
[ 1000 / 1000 ]
```

```
m12.4a_pred <- m12.4a_pred %>%
```

```
  mutate(
    p.pred = apply(link.m12.4a, 2, mean),
    p.link.lo = apply(link.m12.4a, 2, PI)[1,],
    p.link.hi = apply(link.m12.4a, 2, PI)[2,]
  )
```

```
gf_ribbon(p.link.lo + p.link.hi ~ combo + group::"1", data = m12.4a_pred) %>%
  gf_line(p.pred ~ combo + group::"1", data = m12.4a_pred) %>%
  gf_facet_wrap(~ actor)
```



R Code 12.28

```
m12.4a_post <- extract.samples(m12.4a)
m12.4a_postD <- m12.4a_post %>% data.frame()
str(m12.4a_post)
```

```
## List of 6
## $ a_actor      : num [1:16000, 1:7] -0.9048 -1.0586 0.3696 -0.5498 0.0517 ...
## $ a            : num [1:16000(1d)] 0.1379 0.3277 -0.8717 0.0158 0.0439 ...
## $ bp           : num [1:16000(1d)] 0.647 0.916 0.688 0.335 0.329 ...
```

```
## $ bc      : num [1:16000(1d)] -0.505 -0.287 -0.319 -0.382 -0.658 ...
## $ bpc      : num [1:16000(1d)] 0.5761 0.0208 0.3456 0.4342 1.0619 ...
## $ sigma_actor: num [1:16000(1d)] 1.69 1.54 2.39 1.65 4.14 ...
```

```
str(m12.4a_postD)
```


```
## 'data.frame': 16000 obs. of 12 variables:
## $ a_actor.1 : num -0.9048 -1.0586 0.3696 -0.5498 0.0517 ...
## $ a_actor.2 : num 3.37 3.66 5.89 4.36 6.57 ...
## $ a_actor.3 : num -1.151 -1.541 0.176 -0.448 -1.13 ...
## $ a_actor.4 : num -1.311 -1.177 0.237 -0.617 -0.942 ...
## $ a_actor.5 : num -0.8478 -0.6893 0.0311 -0.2395 -0.5979 ...
## $ a_actor.6 : num 0.0258 -0.1118 1.585 0.7382 0.5983 ...
## $ a_actor.7 : num 2.2 1.55 2.63 1.98 2.41 ...
## $ a : num [1:16000(1d)] 0.1379 0.3277 -0.8717 0.0158 0.0439 ...
## $ bp : num [1:16000(1d)] 0.647 0.916 0.688 0.335 0.329 ...
## $ bc : num [1:16000(1d)] -0.505 -0.287 -0.319 -0.382 -0.658 ...
## $ bpc : num [1:16000(1d)] 0.5761 0.0208 0.3456 0.4342 1.0619 ...
## $ sigma_actor: num [1:16000(1d)] 1.69 1.54 2.39 1.65 4.14 ...
```

R Code 12.29

```
gf_dens( ~ a_actor.5, data = m12.4a_postD)
m12.4a_post_long <-
  m12.4a_postD %>%
  tidyr::gather(param, value)
```

```
## Warning: attributes are not identical across measure variables; they will be dropped
```

```
gf_dens( ~ value, data = m12.4a_post_long %>% filter(grepl("actor\\.", param))) %>%
  gf_vline(xintercept = 0, color = "red", alpha = 0.5, linetype = "dashed") %>%
  gf_facet_grid( param ~ .)
```



Creating the link function manually

For some models, it is necessary to create posterior values manually rather than with `link()` or `sim()`. Although this is not required here, we illustrate the manual approach by way of comparison.

R Code 12.30

```
p.link <- function(prosoc_left, condition, actor, post = m12.4a_post) {
  logodds <-
    with(post,
```

```

      a + a_actor[, actor] + (bp + bpc * condition) * prosoc_left)
    return(logistic(logodds))
  }
str(p.link(0, 0, 2))

```

```
## num [1:16000(1d)] 0.971 0.982 0.993 0.988 0.999 ...
```

R Code 12.31

```

G <- expand.grid(
  prosoc_left = 0:1,
  condition = 0:1,
  actor = 1:7) %>%
mutate(
  combo = paste0(prosoc_left, "/", condition)
)

manual_link <-
  with(G, # saves having to type G$ repeatedly
    mapply(p.link, prosoc_left, condition, actor)
  )
# Note: a column here for each row in G
str(manual_link)

```

```
## num [1:16000, 1:28] 0.317 0.325 0.377 0.37 0.524 ...
```

```
dim(G)
```

```
## [1] 28 4
```

```


manual12.4_pred <-
  G %>%
  mutate(
    p.pred = apply(manual_link, 2, mean),
    p.link.lo = apply(manual_link, 2, PI)[1,],
    p.link.hi = apply(manual_link, 2, PI)[2,]
  )

```

```

# replication of plot from above using home spun link
gf_ribbon(p.link.lo + p.link.hi ~ combo + group::"1", data = manual12.4_pred) %>%
  gf_line(p.pred ~ combo + group::"1", data = m12.4a_pred) %>%
  gf_facet_wrap(~ actor)

```



R Code 12.32


```

# don't need multiple actors this time
G <- expand.grid(
  prosoc_left = 0:1,
  condition = 0:1,
  actor = 1) %>%
mutate(
  combo = paste0(prosoc_left, "/", condition)
)

```

R Code 12.33

```

# replace varying intercept samples with zeros
# 1000 samples by 7 actors
a_actor_zeros <- matrix(0, 1000, 7)

```

R Code 12.34

```

# note use of replace list
m12.4a_link2 <-
  link(
    m12.4a,
    n = 1000,
    data = G,
    replace = list(a_actor = a_actor_zeros)
  )

```

```

## [ 100 / 1000 ]
[ 200 / 1000 ]
[ 300 / 1000 ]
[ 400 / 1000 ]
[ 500 / 1000 ]
[ 600 / 1000 ]
[ 700 / 1000 ]
[ 800 / 1000 ]
[ 900 / 1000 ]
[ 1000 / 1000 ]

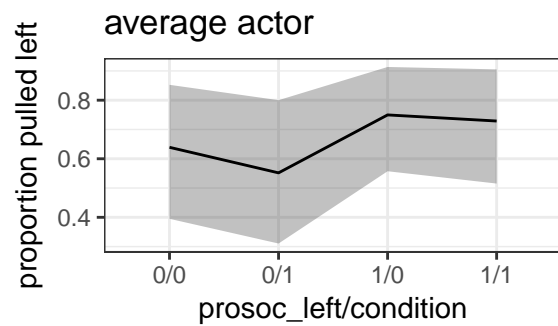
```

```

m12.4a_pred2 <-
  G %>%
  mutate(
    p.pred = apply(m12.4a_link2, 2, mean),
    p.link.lo = apply(m12.4a_link2, 2, PI, prob = 0.8)[1,],
    p.link.hi = apply(m12.4a_link2, 2, PI, prob = 0.8)[2,]
  )

gf_ribbon(p.link.lo + p.link.hi ~ combo + group::"1", data = m12.4a_pred2) %>%
  gf_line(p.pred ~ combo + group::"doesn't matter what this is", data = m12.4a_pred2) %>%
  gf_labs(x = "prosoc_left/condition", y = "proportion pulled left", title = "average actor")

```



R Code 12.35

```
# replace varying intercept samples with simulations
m12.4a_post <- extract.samples(m12.4a)
a_actor_sims <-
  rnorm(7000, 0, m12.4a_post$sigma_actor) %>%
  matrix(1000, 7) # reshape into a 1000 x 7 matrix
```


R Code 12.36

```
m12.4a_link3 <-
  link(
    m12.4a,
    n = 1000,
    data = G,
    replace = list(a_actor = a_actor_sims)
  )
```

```
## [ 100 / 1000 ]
[ 200 / 1000 ]
[ 300 / 1000 ]
[ 400 / 1000 ]
[ 500 / 1000 ]
[ 600 / 1000 ]
[ 700 / 1000 ]
[ 800 / 1000 ]
[ 900 / 1000 ]
[ 1000 / 1000 ]
```

```
m12.4a_pred3 <-
  G %>%
  mutate(
    p.pred = apply(m12.4a_link3, 2, mean),
    p.link.lo = apply(m12.4a_link3, 2, PI, prob = 0.8)[1,],
    p.link.hi = apply(m12.4a_link3, 2, PI, prob = 0.8)[2,]
  )

gf_ribbon(p.link.lo + p.link.hi ~ combo + group::"1", data = m12.4a_pred3) %>%
  gf_line(p.pred ~ combo + group::"doesn't matter what this is", data = m12.4a_pred3) %>%
  gf_labs(x = "prosoc_left/condition", y = "proportion pulled left", title = "marginal of actor")
```



R Code 12.37

```
m12.4a_post <- extract.samples(m12.4a) %>%
  as.data.frame() %>%
  mutate(
    sim_a_actor = rnorm(16000, 0, sigma_actor)
  )

Actors50 <-
  expand.grid(
    actor = 1:50,          # 50 simulated actors
    prosoc_left = 0:1,
    condition = 0:1) %>%
  mutate(
    combo = paste0(prosoc_left, "/", condition),
    logodds =
      m12.4a_post$a[actor] +
      m12.4a_post$sim_a_actor[actor] +
      (m12.4a_post$bp[actor] + m12.4a_post$bpc[actor] * condition) * prosoc_left,
    p = logistic(logodds)
  )
```

R Code 12.38

```
gf_line(p ~ combo + group::actor, data = Actors50, alpha = 0.3) %>%
  gf_labs(
    x = "prosoc_left/condition",
    y = "proportion pulled left",
    title = "50 simulated actors"
  )
```

